

The background features a dark blue field filled with glowing binary code (0s and 1s) in various shades of blue and cyan. Some of the digits are blurred, creating a sense of depth and motion. At the bottom, there are several bright, glowing blue shapes that resemble stylized flames or abstract data points.

GG 606 SCIENTIFIC DATA WRANGLING

Mar 3: Temporal data

TEMPORAL DATA WRANGLING

- Data frequently has a temporal component in addition to a spatial
- dates / times are frequently challenging to work with and need careful consideration when processing
 - if you can, ignore date/time and use integer / arbitrary time units
 - often you cant...

DATE/TIME DATA IN R??

```
data(wind, package = "gstat")  
head(wind[,1:7])
```

##	year	month	day	RPT	VAL	ROS	KIL
## 1	61	1	1	15.04	14.96	13.17	9.29
## 2	61	1	2	14.71	16.88	10.83	6.50
## 3	61	1	3	18.50	16.88	12.33	10.13
## 4	61	1	4	10.58	6.63	11.75	4.58
## 5	61	1	5	13.33	13.25	11.42	6.17
## 6	61	1	6	13.21	8.12	9.96	6.67

```
data(Produc, package = "plm")  
head(Produc[,1:6])
```

##	state	year	pcap	hwy	water	util
## 1	ALABAMA	1970	15032.67	7325.80	1655.68	6051.20
## 2	ALABAMA	1971	15501.94	7525.94	1721.02	6254.98
## 3	ALABAMA	1972	15972.41	7765.42	1764.75	6442.23
## 4	ALABAMA	1973	16406.26	7907.66	1742.41	6756.19
## 5	ALABAMA	1974	16762.67	8025.52	1734.85	7002.29
## 6	ALABAMA	1975	17316.26	8158.23	1752.27	7405.76

```
library(foreign)  
read.dbf(system.file("shapes/sids.dbf", package="maptools"))[,1:5,c(5,9:14)]
```

##	NAME	BIR74	SID74	NWBIR74	BIR79	SID79	NWBIR79
## 1	Ashe	1091	1	10	1364	0	19
## 2	Alleghany	487	0	10	542	3	12
## 3	Surry	3188	5	208	3616	6	260
## 4	Currituck	508	1	123	830	2	145
## 5	Northampton	1421	9	1066	1606	3	1197

lubridate basics

- There are three types of date/time data that refer to an instant in time:
 - A **date**. Tibbles print this as <date>.
 - A **time** within a day. Tibbles print this as <time>.
 - no need to handle this generally, can use a double of integer generally
 - A **date-time** is a date plus a time: it uniquely identifies an instant in time (typically to the nearest second). Tibbles print this as <dtm>.
 - base R uses a class called POSIXct

lubridate basics

- Sometimes we need to create dates based on the current execution time in R
- usually however we are reading in date and/or date-time data as character data in data files
- if we are connected to a database or have a database file with actual data types (e.g., dbf files) we can read directly into date type classes in R
- date data type is actually a double

```
> Sys.Date()
[1] "2022-03-02"
> class(Sys.Date())
[1] "Date"
> library(lubridate)
> today()
[1] "2022-03-02"
> class(today())
[1] "Date"
> now()
[1] "2022-03-02 12:36:06 PST"
> class(now())
[1] "POSIXct" "POSIXt"
> typeof(today())
[1] "double"
> typeof(Sys.Date())
[1] "double"
```

lubridate - input

- Dates from character data
- identify the order in which year, month, and day appear in your dates, then arrange “y”, “m”, and “d” in the same order
 - gives you the name of the lubridate function that will parse your date

lubridate – dates + times

```
ymd_hms("2017-01-31 20:11:59")
#> [1] "2017-01-31 20:11:59 UTC"
mdy_hm("01/31/2017 08:01")
#> [1] "2017-01-31 08:01:00 UTC"
```

lubridate - dates

```
ymd("2017-01-31")
#> [1] "2017-01-31"

mdy("January 31st, 2017")
#> [1] "2017-01-31"
dmy("31-Jan-2017")
#> [1] "2017-01-31"
```

base r

```
as.Date("2017-01-31", "%Y-%m-%d")
[1] "2017-01-31"
> as.Date("January 31, 2017", "%B %d, %Y")
[1] "2017-01-31"
as.Date("31-Jan-2017", "%d-%b-%Y")
[1] "2017-01-31"
```

lubridate - input

Data frame with columns

year, month, day

Date of departure.

dep_time, arr_time

Actual departure and arrival times (format HHMM or HMM), local tz.

sched_dep_time, sched_arr_time

Scheduled departure and arrival times (format HHMM or HMM), local tz.

dep_delay, arr_delay

Departure and arrival delays, in minutes. Negative times represent early departures/arrivals.

- Equally common – because dates are such as pain – is that date data are split across columns, making it easy to query and filter using integer data types
- to filter by month or by year we can use standard filtering techniques and not have to use date classes at all
- to explore arrival and departure times we must create date-time data using `make_datetime`
 - assemble from existing columns

```
flights %>%
  select(year, month, day, hour, minute) %>%
  mutate(departure = make_datetime(year, month, day,
    hour, minute))
#> # A tibble: 336,776 x 6
#>   year month   day hour minute departure
#>   <int> <int> <int> <dbl> <dbl> <dtm>
#> 1  2013     1     1     5     15 2013-01-01 05:15:00
#> 2  2013     1     1     5     29 2013-01-01 05:29:00
#> 3  2013     1     1     5     40 2013-01-01 05:40:00
#> 4  2013     1     1     5     45 2013-01-01 05:45:00
#> 5  2013     1     1     6     0  2013-01-01 06:00:00
#> 6  2013     1     1     5     58 2013-01-01 05:58:00
#> # ... with 336,770 more rows
```

lubridate — extracting parts

- pull out individual parts of the date with the accessor functions
 - `year()`, `month()`
 - `mday()` (day of the month)
 - `yday()` (day of the year)
 - `wday()` (day of the week)
 - `hour()`, `minute()`, and `second()`

```
datetime <- ymd_hms("2016-07-08 12:34:56")

year(datetime)
#> [1] 2016
month(datetime)
#> [1] 7
month(datetime, label = TRUE)
#> [1] Jul
mday(datetime)
#> [1] 8
yday(datetime)
#> [1] 190
wday(datetime)
#> [1] 6
wday(datetime, label = TRUE, abbr = FALSE)
#> [1] Friday
```


lubridate – manipulating date data

- Often you need to do basic arithmetic with date data:
 - e.g., calculate the duration between a start date end date or start time and end time
- `difftime` class object records a time span of seconds, minutes, hours, days, or weeks.
- `as.duration` provides a response always in seconds
- constructors such as `ddays` and `dyears` can be used for arithmetic operations

```
> Sys.Date() - 4
[1] "2022-02-26"
> class(Sys.Date() - 4)
[1] "Date"
> Sys.Date() - (Sys.Date() - 4)
Time difference of 4 days
> class(Sys.Date() - (Sys.Date() - 4))
[1] "difftime"
> as.duration(Sys.Date() - (Sys.Date() - 4))
[1] "345600s (~4 days)"
```

```
tomorrow <- today() + ddays(1)
last_year <- today() - dyears(1)
```

lubridate – manipulating date data

- Some planes appear to have arrived at their destination *before* they departed from New York City.

```
flights_dt %>%
  filter(arr_time < dep_time)
#> # A tibble: 10,633 x 9
#>   origin dest  dep_delay arr_delay dep_time          sched_dep_time
#>   <chr>  <chr>    <dbl>    <dbl> <dtm>          <dtm>
#> 1 EWR    BQN         9        -4 2013-01-01 19:29:00 2013-01-01 19:20:00
#> 2 JFK    DFW        59         NA 2013-01-01 19:39:00 2013-01-01 18:40:00
#> 3 EWR    TPA        -2         9 2013-01-01 20:58:00 2013-01-01 21:00:00
#> 4 EWR    SJU        -6       -12 2013-01-01 21:02:00 2013-01-01 21:08:00
#> 5 EWR    SFO        11       -14 2013-01-01 21:08:00 2013-01-01 20:57:00
#> 6 LGA    FLL       -10        -2 2013-01-01 21:20:00 2013-01-01 21:30:00
#> # ... with 10,627 more rows, and 3 more variables: arr_time <dtm>,
#> #   sched_arr_time <dtm>, air_time <dbl>
```

lubridate – manipulating date data

- We can fix this by adding `days(1)` to the arrival time of each overnight flight

```
flights_dt <- flights_dt %>%  
  mutate(  
    overnight = arr_time < dep_time,  
    arr_time = arr_time + days(overnight * 1),  
    sched_arr_time = sched_arr_time + days(overnight * 1)  
  )
```

```
> head(flights_dt)  
# A tibble: 6 × 10  
  origin dest dep_delay arr_delay dep_time sched_dep_time arr_time sched_arr_time air_time overnight  
  <chr>  <chr>    <dbl>    <dbl> <dtm>          <dtm>          <dtm>          <dtm>          <dbl>  <lgl>  
1 EWR    IAH         2        11 2013-01-01 05:17:00 2013-01-01 05:15:00 2013-01-01 08:30:00 2013-01-01 08:19:00    227 FALSE  
2 LGA    IAH         4        20 2013-01-01 05:33:00 2013-01-01 05:29:00 2013-01-01 08:50:00 2013-01-01 08:30:00    227 FALSE  
3 JFK    MIA         2        33 2013-01-01 05:42:00 2013-01-01 05:40:00 2013-01-01 09:23:00 2013-01-01 08:50:00    160 FALSE  
4 JFK    BQN        -1       -18 2013-01-01 05:44:00 2013-01-01 05:45:00 2013-01-01 10:04:00 2013-01-01 10:22:00    183 FALSE  
5 LGA    ATL        -6       -25 2013-01-01 05:54:00 2013-01-01 06:00:00 2013-01-01 08:12:00 2013-01-01 08:37:00    116 FALSE  
6 EWR    ORD        -4        12 2013-01-01 05:54:00 2013-01-01 05:58:00 2013-01-01 07:40:00 2013-01-01 07:28:00    150 FALSE
```


LUBRIDATE — MANIPULATING DATE DATA

```
# days of the week:
dow = c("Sunday", "Monday", "Tuesday",
"Wednesday", "Thursday", "Friday", "Saturday")

flights_dow <- flights_dt %>%
  mutate(hour = lubridate::hour(sched_dep_time))
  %>%


  mutate(day = factor(weekdays(sched_dep_time),
levels = dow))

flights_dow %>%
  filter(!is.na(hour) & !day %in% c("Saturday",
"Sunday")) %>%
  ggplot(aes(hour)) +
    geom_bar() +
    facet_wrap(~day, nrow = 1)
```




ANALYSIS OF TEMPORAL DATA

- Time series analysis – lots of options available in
- a time series is a set of observations taken sequentially in time
- We often want to plot time series
 - with time increasing almost always on the x-axis
- For time series we typically want arbitrary time units
 - i.e., data stored as integers



ANALYSIS OF TEMPORAL DATA

```
data(lynx, package = "datasets")  
plot.ts(lynx, ylab = "# of Lynx trapped in Canada", las = 1, col = "blue", lwd = 2)
```



ANALYSIS OF TEMPORAL DATA

- Common forms of time series analysis available in R
 - time series modelling
 - time series decomposition
 - trend, seasonal, error components
 - Box Jenkins modelling
 - ARMA – autoregressive moving average (forecast package)
 - ARIMA – autoregressive integrated moving average (forecast package)
 - Random walks, correlated random walks, Brownian movement, Levy models
 - Change point analysis
 - significance shifts in mean, variance, etc.

<https://cran.uni-muenster.de/web/views/TimeSeries.html>

xts AND zoo PACKAGES

- For time series analysis while maintaining date/time data classes
-

Introducing xts and zoo objects

MANIPULATING TIME SERIES DATA WITH XTS AND ZOO IN R



| spacetime

SUMMARY

- There are many different tools available for handling temporal data in R
 - ts objects
 - date classes and lubridate
 - zoo and xts packages
- In general, select the simplest packages/classes you can for your needs and that integrate with your overall workflow
- Build on your knowledge of spatial and temporal classes to explore space/time data





GG 606 SCIENTIFIC DATA WRANGLING

1. Jan 27: databases